

Project: Traffic Sign Recognition

Build a Traffic Sign Recognition Project

The goals / steps of this project are the following:

- Load the data set (see below for links to the project data set)
- Explore, summarize and visualize the data set
- Design, train and test a model architecture
- Use the model to make predictions on new images
- Analyze the softmax probabilities of the new images
- Summarize the results with a written report

Rubric Points

Here I will consider the [rubric points](#) individually and describe how I addressed each point in my implementation.

Data Set Summary & Exploration


1. Basic summary of the data set.

Once the data sets are imported, I used the numpy library to calculate summary statistics of the traffic signs data set:

- The size of training set is 3799
- The size of the validation set is 4410
- The size of test set is 12630
- The shape of a traffic sign image is (32, 32, 3)
- The number of unique classes/labels in the data set is 43

2. Exploratory visualization of the dataset.

Here is a horizontal histogram chart showing how the three data sets are distributed along the different labels.



In order to have a better idea of the training set, for each label, I printed out 4 random images. The output is visible in the image below.

0 : Speed limit (20km/h)				
1 : Speed limit (30km/h)				
2 : Speed limit (50km/h)				
3 : Speed limit (60km/h)				
4 : Speed limit (70km/h)				
5 : Speed limit (80km/h)				
6 : End of speed limit (80km/h)				
7 : Speed limit (100km/h)				
8 : Speed limit (120km/h)				
9 : No passing				
10 : No passing for vehicles over 3.5 metric tons				
11 : Right-of-way at the next intersection				
12 : Priority road				
13 : Yield				
14 : Stop				
15 : No vehicles				
16 : Vehicles over 3.5 metric tons prohibited				
17 : No entry				
18 : General caution				
19 : Dangerous curve to the left				
20 : Dangerous curve to the right				

21 : Double curve



22 : Bumpy road



23 : Slippery road



24 : Road narrows on the right



25 : Road work



26 : Traffic signals



27 : Pedestrians



28 : Children crossing



29 : Bicycles crossing



30 : Beware of ice/snow



31 : Wild animals crossing



32 : End of all speed and passing limits



33 : Turn right ahead



34 : Turn left ahead



35 : Ahead only



36 : Go straight or right



37 : Go straight or left



38 : Keep right



39 : Keep left



40 : Roundabout mandatory



41 : End of no passing



42 : End of no passing by vehicles over 3.5 metric tons



Design and Test a Model Architecture

2. Image data preprocess

As a first step, I converted the image to grayscale. I decided to grayscale them by simply averaging the Red Green and Blue values of the original image.

After, I normalized the images in order to have the grayscale values inside the range $[-1, 1]$.

Keeping the gray values inside this range instead of $[0, 255]$ is fundamental for the training operation in order to get good performances of the model.

Then I shuffled the data using the specific function.

Note: in the first project submission, I decided to omit the grayscaling process since in terms of accuracy it does not give us better result.

Thanks to the reviewer I realized that this method helps to drastically reduce the image dimension and the computational load. In fact, keeping approximately the same accuracy, the images to be processed have 1/3 of the original RGB image dimension.

2. Model architecture

My final model consisted of a standard LeNet architecture and, in particular, is made of the following layers:

Layer	Description
Input	32x32x3 RGB image
Convolution 1	1x1 stride, VALID padding, output: 28x28x6
RELU	
Max pooling 1	2x2 stride, VALID padding, output: 14x14x6
Convolution 2	1x1 stride, VALID padding, output: 10x10x16
RELU	
Max pooling 2	2x2 stride, VALID padding, output: 5x5x16
Flatten	Output: 400
Fully connected	input: 400, output: 120
RELU	
Fully connected	input: 120, output: 84
RELU	
Fully connected	input: 84, output: 43

3. Model training

To train the model, after some tuning, I decided to use the following hyperparameters:

- Epochs = 20
- Batch size = 128
- Learning rate = 0.001

For the training optimization, I used the `softmax_cross_entropy_with_logits` function which compares the evaluated logits with the one hot end labels.

This gives the cross entropy which is a measure of how different the logits are from the ground truth training labels.

Next, the cross-entropy values are averaged along all the training images.

Then, I minimize the loss function using the `AdamOptimizer` function which is similar to the stochastic gradient descent method.

4. Results

My final model results were:

- training set accuracy of 0.993
- validation set accuracy of 0.976
- test set accuracy of 0.891

These results satisfy the minimum requests for the project completion.

Test a Model on New Images

1. Chosen images

Here are five German traffic signs that I found on the web:



In these images, I picked signs that might have some additional difficulty in the prediction.

For example, the third image 'Go straight or left' is taken from a steep angle.

In addition, the fourth image 'Road work' presents some differences from the one in the dataset. In fact, the background is yellow instead of white.

2. Model's predictions on the new traffic signs.

Here are the results of the prediction on the new images:

Image		Prediction
1	Yield	Yield
2	Dangerous curve to the right	Dangerous curve to the right
3	Go straight or left	Go straight or left
4	Road work	Road work
5	Road narrows to the right	Road narrows to the right

The model was able to correctly guess all the 5 traffic signs, which gives an accuracy of 100%. This result is compatible with the one obtained evaluating the accuracy on the given training set.

3. Softmax probabilities

Analyzing the top five softmax probabilities for each prediction allows to understand how the model works. For most of the images, the model recognizes with absolute precision the traffic signs.

In the last two images, the correct prediction does not have 100% probability but it is still close to it.

Image 1: Yield	
Probability	Prediction
100%	Yield
0%	Ahead only
0%	No vehicles
0%	No passing
0%	Speed limit (50km/h)

Image 2: Dangerous curve to the right	
Probability	Prediction
99.87%	Dangerous curve to the right
0.10%	Slippery road
0%	Road work
0%	Right-of-way at the next intersection
0%	Children crossing

Image 3 : Go straight or left	
Probability	Prediction
100%	Go straight or left
0%	Dangerous curve to the left
0%	Keep left
0%	No passing for vehicles over 3.5 metric tons
0%	Traffic signals

Image 4: Road work	
Probability	Prediction
51.42%	Road work
47.19%	Priority road
1.39%	Right-of-way at the next intersection
0%	Speed limit (50km/h)
0%	Speed limit (80km/h)

Image 5: Road narrow to the right	
Probability	Prediction
93.85%	Road narrows to the right
5.81%	Bicycles crossing
0.23%	Speed limit (50km/h)
0.07%	Children crossing
0.03%	Road work

Conclusions

The trained model has a high prediction accuracy and it is also robust since it correctly works with other images taken from internet.

The overall result might be improved by applying some changes:

- use more sophisticate architectures
- apply data augmentation (rotation, translation, zoom, flips, and/or color perturbation) in order to have a better view of the signs
- use more layers on the LeNet architecture

